

# Foundations of Machine Learning

## Module 2: Linear Regression and Decision Tree

### Part D: Overfitting


Sudeshna Sarkar  
IIT Kharagpur

# Overfitting


- Learning a tree that classifies the training data perfectly may not lead to the tree with the best generalization performance.
  - There may be noise in the training data
  - May be based on insufficient data
- A hypothesis  $h$  is said to overfit the training data if there is another hypothesis,  $h'$ , such that  $h$  has smaller error than  $h'$  on the training data but  $h$  has larger error on the test data than  $h'$ .

# Overfitting


- Learning a tree that classifies the training data perfectly may not lead to the tree with the **best generalization performance**.
  - There may be noise in the training data
  - May be based on insufficient data
- A hypothesis  $h$  is said to **overfit the training data** if there is another hypothesis,  $h'$ , such that  $h$  has smaller error than  $h'$  on the training data but  $h$  has larger error on the test data than  $h'$ .



# Overfitting



# Underfitting and Overfitting (Example)



500 circular and 500 triangular data points.

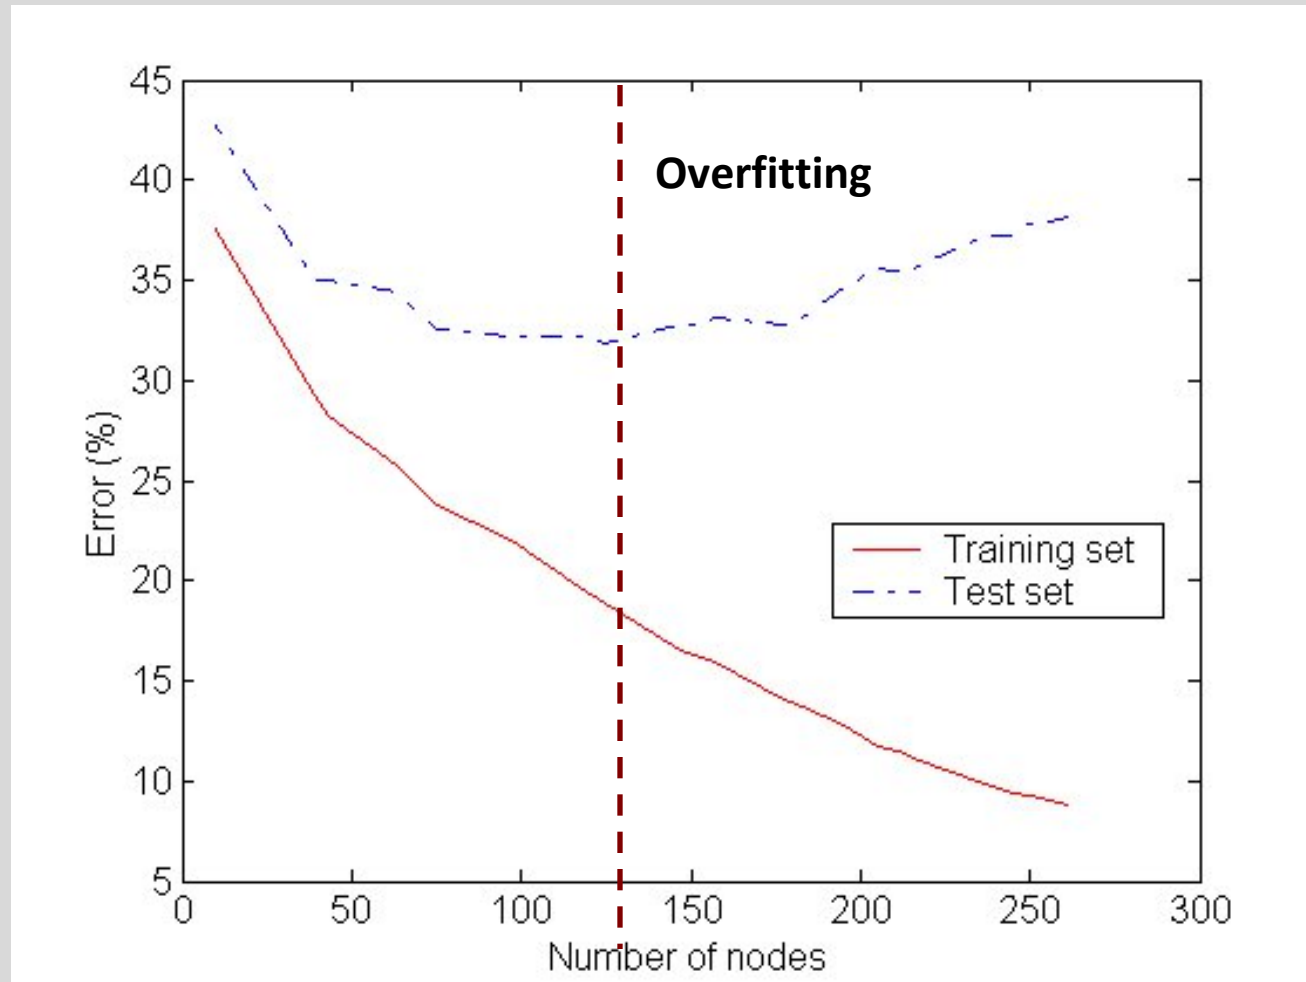
Circular points:

$$0.5 \bullet \sqrt{x_1^2 + x_2^2} \bullet 1$$

Triangular points:


$$\sqrt{x_1^2 + x_2^2} > 0.5 \text{ or } \sqrt{x_1^2 + x_2^2} < 1$$

# Underfitting and Overfitting



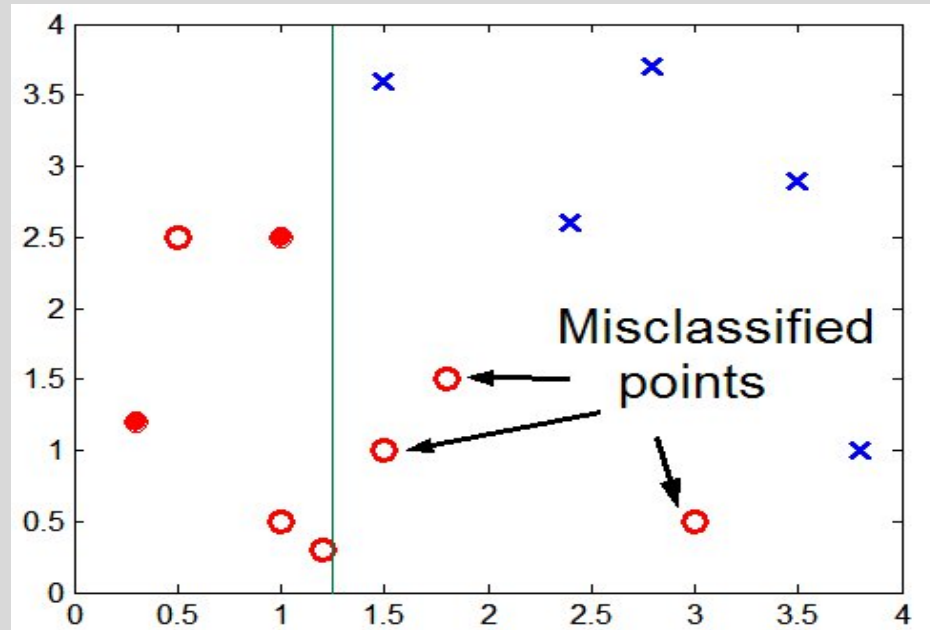
**Underfitting:** when model is too simple, both training and test errors are large

# Overfitting due to Noise



Decision boundary is distorted by noise point

# Overfitting due to Insufficient Examples



Lack of data points makes it difficult to predict correctly the class labels of that region



# Notes on Overfitting

- Overfitting results in decision trees that are more complex than necessary
- Training error no longer provides a good estimate of how well the tree will perform on previously unseen records

# Avoid Overfitting

- How can we avoid overfitting a decision tree?
  - **Prepruning**: Stop growing when data split not statistically significant
  - **Postpruning**: Grow full tree then remove nodes
- Methods for evaluating subtrees to prune:
  - Minimum description length (MDL):  
Minimize:  $\text{size}(\text{tree}) + \text{size}(\text{misclassifications}(\text{tree}))$
  - Cross-validation

# Pre-Pruning (Early Stopping)

- Evaluate splits before installing them:
  - Don't install splits that don't look worthwhile
  - when no worthwhile splits to install, done

# Pre-Pruning (Early Stopping)


- Typical stopping conditions for a node:
  - Stop if all instances belong to the same class
  - Stop if all the attribute values are the same
- More restrictive conditions:
  - Stop if number of instances is less than some user-specified threshold
  - Stop if class distribution of instances are independent of the available features (e.g., using  $\chi^2$  test)
  - Stop if expanding the current node does not improve impurity measures (e.g., Gini or information gain).

# Reduced-error Pruning

- A post-pruning, cross validation approach
  - Partition training data into “grow” set and “validation” set.
  - Build a complete tree for the “grow” data
  - Until accuracy on validation set decreases, do:
    - For each non-leaf node in the tree
    - Temporarily prune the tree below; replace it by majority vote
    - Test the accuracy of the hypothesis on the validation set
    - Permanently prune the node with the greatest increase in accuracy on the validation test.
- Problem: Uses less data to construct the tree
- Sometimes done at the rules level

General Strategy: Overfit and Simplify

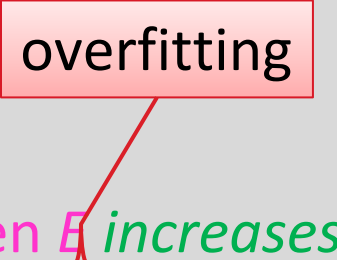
# Reduced Error Pruning



# Model Selection & Generalization

- Learning is an **ill-posed problem**; data is not sufficient to find a unique solution
- The need for **inductive bias**, assumptions about  $H$
- **Generalization**: How well a model performs on new data
- Overfitting:  $H$  more complex than  $C$  or  $f$
- Underfitting:  $H$  less complex than  $C$  or  $f$

# Triple Trade-Off

- There is a trade-off between three factors:
    - Complexity of  $H$ ,  $c(H)$ ,
    - Training set size,  $N$ ,
    - Generalization error,  $E$  on new data
  - As  $N$  *increases*,  $E$  *decreases*
  - As  $c(H)$  *increases*, first  $E$  *decreases* and then  $E$  *increases*
  - As  $c(H)$  *increases*, the training error *decreases* for some time and then stays constant (frequently at 0)
- 



# Notes on Overfitting

- **overfitting** happens when a model is capturing idiosyncrasies of the data rather than generalities.
  - Often caused by too many parameters relative to the amount of training data.
  - E.g. an order- $N$  polynomial can intersect any  $N+1$  data points

# Dealing with Overfitting

- Use more data
- Use a tuning set
- **Regularization**
- Be a Bayesian

# Regularization

- In a linear regression model overfitting is characterized by large weights.

|       | $M = 0$ | $M = 1$ | $M = 3$ | $M = 9$     |
|-------|---------|---------|---------|-------------|
| $w_0$ | 0.19    | 0.82    | 0.31    | 0.35        |
| $w_1$ |         | -1.27   | 7.99    | 232.37      |
| $w_2$ |         |         | -25.43  | -5321.83    |
| $w_3$ |         |         | 17.37   | 48568.31    |
| $w_4$ |         |         |         | -231639.30  |
| $w_5$ |         |         |         | 640042.26   |
| $w_6$ |         |         |         | -1061800.52 |
| $w_7$ |         |         |         | 1042400.18  |
| $w_8$ |         |         |         | -557682.99  |
| $w_9$ |         |         |         | 125201.43   |

# Penalize large weights in Linear Regression

- Introduce a penalty term in the loss function.

$$E(\vec{w}) = \frac{1}{2} \sum_{n=0}^{N-1} \{t_n - y(x_n, \vec{w})\}^2$$

## Regularized Regression

1. (L2-Regularization or Ridge Regression)

$$E(\vec{w}) = \frac{1}{2} \sum_{n=0}^{N-1} (t_n - y(x_n, \vec{w}))^2 + \frac{\lambda}{2} \|\vec{w}\|^2$$

1. L1-Regularization

$$E(\vec{w}) = \frac{1}{2} \sum_{n=0}^{N-1} (t_n - y(x_n, \vec{w}))^2 + \lambda |\vec{w}|_1$$