

Foundations of Machine Learning




Module 6: Neural Network

Part C: Neural Network and Backpropagation Algorithm


Sudeshna Sarkar
IIT Kharagpur

Single layer Perceptron

- Single layer perceptrons learn linear decision boundaries





x: class I ($y = 1$)
o: class II ($y = -1$)



x: class I ($y = 1$)
o: class II ($y = -1$)



Boolean OR

input x_1	input x_2	ouput
0	0	0
0	1	1
1	0	1
1	1	1




Boolean AND

input x_1	input x_2	ouput
0	0	0
0	1	0
1	0	0
1	1	1




Boolean XOR

input x1	input x2	ouput
0	0	0
0	1	1
1	0	1
1	1	0



Boolean XOR


input x1	input x2	ouput
0	0	0
0	1	1
1	0	1
1	1	0



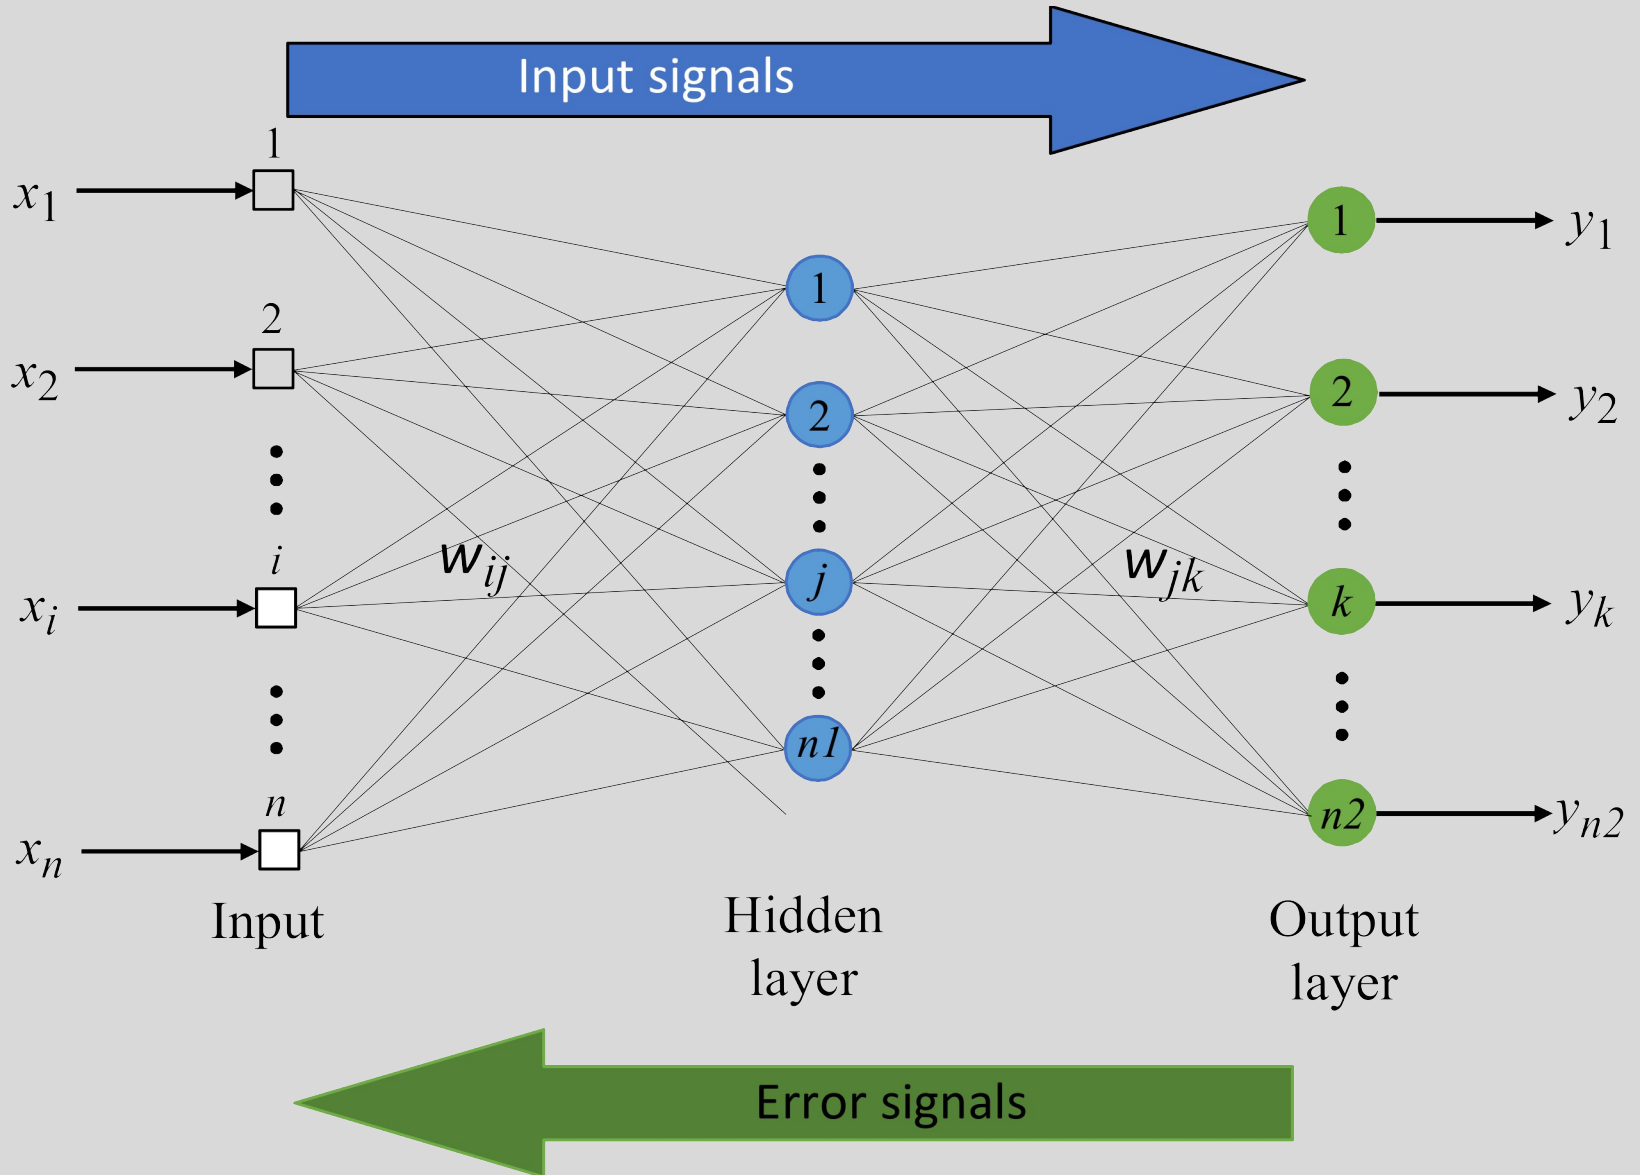
Representation Capability of NNs

- Single layer nets have limited representation power (linear separability problem). Multi-layer nets (or nets with non-linear hidden units) may overcome linear inseparability problem.
- Every Boolean function can be represented by a network with a single hidden layer.
- Every bounded continuous function can be approximated with arbitrarily small error, by network with one hidden layer
- Any function can be approximated to arbitrary accuracy by a network with two hidden layers.

Multilayer Network



Two-layer back-propagation neural network



Derivation

- For one output neuron, the error function is

$$E = \frac{1}{2} (y - o)^2$$

- For each unit j , the output o_j is defined as

$$o_j = \varphi(\text{net}_j) = \varphi\left(\sum_{k=1}^n w_{kj} o_k\right)$$

The input net_j to a neuron is the weighted sum of outputs o_k of previous n neurons.

- Finding the derivative of the error:

$$\frac{\partial E}{\partial w_{ij}} = \frac{\partial E}{\partial o_j} \frac{\partial o_j}{\partial \text{net}_j} \frac{\partial \text{net}_j}{\partial w_{ij}}$$

For one output neuron, the error function is $E = \frac{1}{2} (y - o)^2$

For each unit j , the output o_j is defined as

$$o_j = \varphi(\text{net}_j) = \varphi\left(\sum_{k=1}^n w_{kj} o_k\right)$$

$$\begin{aligned} \frac{\partial E}{\partial w_{ij}} &= \frac{\partial E}{\partial o_j} \frac{\partial o_j}{\partial \text{net}_j} \frac{\partial \text{net}_j}{\partial w_{ij}} \\ &= \sum_l \left(\frac{\partial E}{\partial o_l} \frac{\partial o_l}{\partial \text{net}_{z_l}} w_{jl} \right) \varphi(\text{net}_j) (1 - \varphi(\text{net}_j)) o_i \\ \frac{\partial E}{\partial w_{ij}} &= \delta_j o_i \end{aligned}$$

with

$$\delta_j = \frac{\partial E}{\partial o_j} \frac{\partial o_j}{\partial \text{net}_j} = \begin{cases} (o_j - y_j) o_j (1 - o_j) & \text{if } j \text{ is an output neuron} \\ \left(\sum_z \delta_{z_l} w_{jl} \right) o_j (1 - o_j) & \text{if } j \text{ is an inner neuron} \end{cases}$$

To update the weight w_{ij} using gradient descent, one must choose a learning rate η .

$$\Delta w_{ij} = -\eta \frac{\partial E}{\partial w_{ij}}$$

Backpropagation Algorithm

Initialize all weights to small random numbers.

Until satisfied, do

– For each training example, do

- Input the training example to the network and compute the network outputs

- For each output unit k

$$\delta_k \leftarrow o_k(1-o_k)(y_k-o_k)$$

- For each hidden unit h

$$\delta_h \leftarrow o_h(1-o_h)\sum_{k \in \text{outputs}} w_{h,k}\delta_k,$$

- Update each network weight $w_{i,j}$

$$w_{i,j} \leftarrow w_{i,j} + \Delta w_{i,j}$$

where

$$\Delta w_{i,j} = \eta \delta_j x_{i,j}$$

x_d = input

y_d = target output

o_d = observed unit output

w_{ij} = wt from i to j

Backpropagation

- Gradient descent over entire network weight vector
- Can be generalized to arbitrary directed graphs
- Will find a local, not necessarily global error minimum
- May include weight momentum α

$$\Delta w_{i,j}(n) = \eta \delta_j x_{i,j} + \alpha \Delta w_{i,j}(n-1)$$

- Training may be slow.
- Using network after training is very fast

Training practices: batch vs. stochastic vs. mini-batch gradient descent

- **Batch gradient descent:**
 1. Calculate outputs for the entire dataset
 2. Accumulate the errors, back-propagate and update
- **Stochastic/online gradient descent:**
 1. Feed forward a training example
 2. Back-propagate the error and update the parameters
- **Mini-batch gradient descent:**

Too slow to converge
Gets stuck in local minima

Converges to the solution faster
Often helps get the system out of
local minima

Learning in *epochs*


Stopping

- Train the NN on the entire training set over and over again
- Each such episode of training is called an “epoch”



Stopping

1. Fixed maximum number of epochs: most naïve
2. Keep track of the training and validation error curves.

Overfitting in ANNs



Local Minima



- NN can get stuck in local minima for small networks.
- For most large networks (many weights) local minima rarely occurs.
- It is unlikely that you are in a minima in every dimension simultaneously.

ANN

- Highly expressive non-linear functions
- Highly parallel network of logistic function units
- Minimizes sum of squared training errors
- Can add a regularization term (weight squared)
- Local minima
- Overfitting

Thank You